Hritvik Dinkar

Bengaluru, India | +91 9876842631 | hritvik.works@gmail.com
linkedin.com/in/hritvik | hritvik-dinkar.pages.dev

Senior Backend / Platform Engineer | Generative AI & Distributed Systems

Software Development Engineer III at Meesho with 8+ years building production backend, GenAI, and data platforms. Owns system design, implementation, rollout, and operations for supplier-support automation and shared LLM services, with a focus on fault tolerance, configurable orchestration, and production reliability.

TECHNICAL SKILLS

Generative AI: Azure OpenAI, LLM inference routing, provider/deployment failover, prompt templates and registries, model/deployment registries, deterministic traffic splitting, LLM evaluation, intent classification, response generation, agent loops, FAQ retrieval

System Design: Distributed systems, microservices, event-driven architecture, API composition, dependency modeling, concurrent execution, fan-out/fan-in, idempotency, retries, dead-letter queues (DLQs), replay, circuit breakers, caching, high availability, capacity planning

Languages & Platforms: Java, Python, Spring Boot, FastAPI, REST APIs, Kafka, MongoDB, Redis, Elasticsearch, Apache Spark, Prometheus, Grafana, Metabase, Terraform, Jenkins

EXPERIENCE

Meesho | Software Development Engineer III (promoted from SDE II)

May 2022 - Present | Bengaluru, India

- Owned end-to-end architecture and delivery of Agentless, a supplier-support automation platform combining Java/Spring Boot workflow orchestration with Python/FastAPI GenAI services; expanded automation to 60% of supplier ticket categories.
- Designed durable, event-driven workflows using Kafka fan-out/fan-in, MongoDB ticket/entity processing contexts, step-level idempotency, typed retries, DLQs, replay, and stalled-ticket recovery to resume partial failures safely.
- Converted business-owned SOPs into configurable workflow steps and MVEL rules, integrating LLM-based intent classification, response generation, and human escalation paths.
- Built shared GenAI platform services with model/deployment and prompt registries, a provider-agnostic LLM bridge, validated routing plans, deterministic weighted traffic splitting, and automatic provider/deployment failover.
- Implemented LLM evaluation flows and Prometheus/Grafana observability; supported configuration refresh and Redis caching, with Terraform-backed Azure OpenAI deployment operations.
- Designed a configuration-driven API-composition service with field/API dependency modeling, concurrent execution, fan-out, and response stitching; guided a junior engineer through execution-layer implementation and configuration-only onboarding.
- Reduced earlier Support Chatbot human escalation from 25% to about 10% (low of 8%) through FAQ retrieval, model routing, prompt and evaluation improvements, fallback policy, and agent-loop response generation using ticket/order API context.
- Validated support systems at 5x projected sale traffic and helped resolve a thread-pool memory leak, reducing pod requirements from 35 cores / 70 GB to 21 cores / 40 GB.
- Cut contact-center data ingestion runtime from 2 hours to 2 minutes; delivered zero-downtime ticketing migration and reconciliation workflows across external support platforms.
- Mentored engineers through task decomposition, architecture and code reviews; coordinated rollout and stakeholder demos across product, support, and business teams.

Oracle | Software Developer

Jun 2021 - May 2022 | Hyderabad, India

- Built an embeddable Oracle Commerce integration plugin for credential collection, field mapping, sync initiation, and sync monitoring using Oracle Visual Builder, JavaScript, and Knockout.

Reliance Jio | Software Developer

Jul 2018 - May 2021 | Navi Mumbai, India

- Built Java/Python computation and analytics services using Elasticsearch, Kafka, and Spark, supporting KPI processing, forecasting, dashboards, and high-availability cluster operations.
- Implemented dependency-aware derived-KPI computation with topological evaluation and cached multi-resolution aggregations for large-scale telecom analytics.

EDUCATION

Indian Institute of Technology, Ropar | B.Tech., Electrical and Electronics Engineering | 2014 - 2018